

COMP 445
Data Communications & Computer networks
Winter 2022

Application Layer

- ✓ Principles of network applications
- ✓ Web and HTTP
- ✓ Electronic mail
- ✓ DNS
- ✓ P2P applications
- ✓ Video streaming and CDN
- ✓ Sockets

Application Layer – Part 3

- ✓ P2P Applications
 - ✓ Characteristics
 - ✓ File distribution with BitTorrent
- ✓ Video streaming and Content Delivery Networks
 - ✓ Internet video
 - ✓ HTTP Streaming
 - ✓ CDN

Learning objectives

- To quantify the differences between file distribution using client-server vs. P2P architectures
- To describe the operation of BitTorrent as on P2P application
- To explain how the video streaming services are implemented and the application-layer protocols involved
- To describe the way multimedia content can be distributed using Content Delivery Networks.

Application Layer – Part 3


- ✓ P2P Applications
 - ✓ Characteristics
 - ✓ File distribution with BitTorrent
- ✓ Video streaming and Content Delivery Networks
 - ✓ Internet video
 - ✓ HTTP Streaming
 - ✓ CDN

Pure P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:

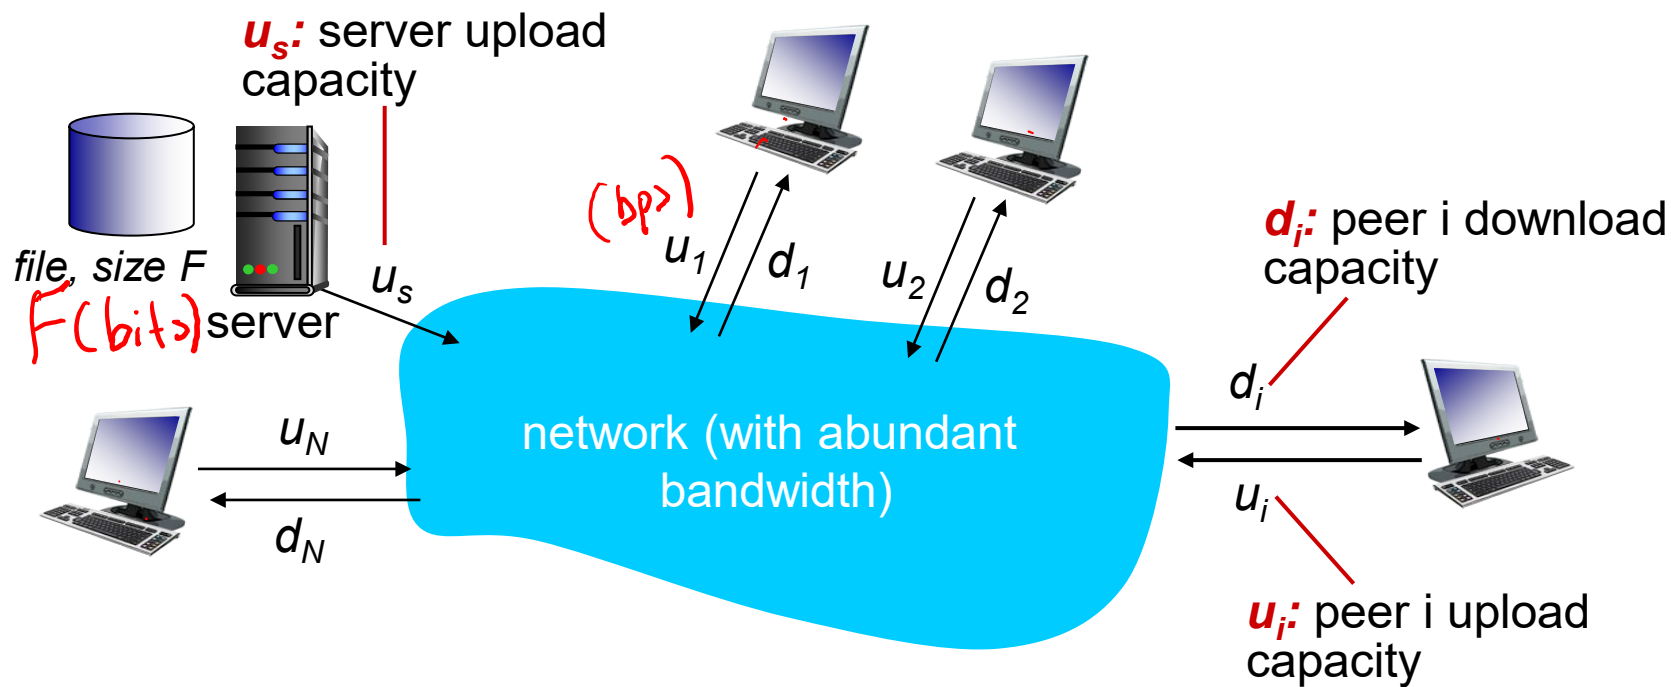
- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



File distribution: client-server vs P2P

Question: how much time to distribute file (size F) from one server to N peers?


- peer upload/download capacity is limited resource



File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:

- time to send one copy: F/u_s
- time to send N copies: NF/u_s



- **client:** each client must download file copy

- d_{\min} = min client download rate
- min client download time: F/d_{\min}

access network
↓
bottleneck


time to distribute F
to N clients using
client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N

File distribution time: P2P

- **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- **client:** each client must download file copy
 - min client download time: F/d_{\min}
- **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$



time to distribute F
to N clients using
P2P approach

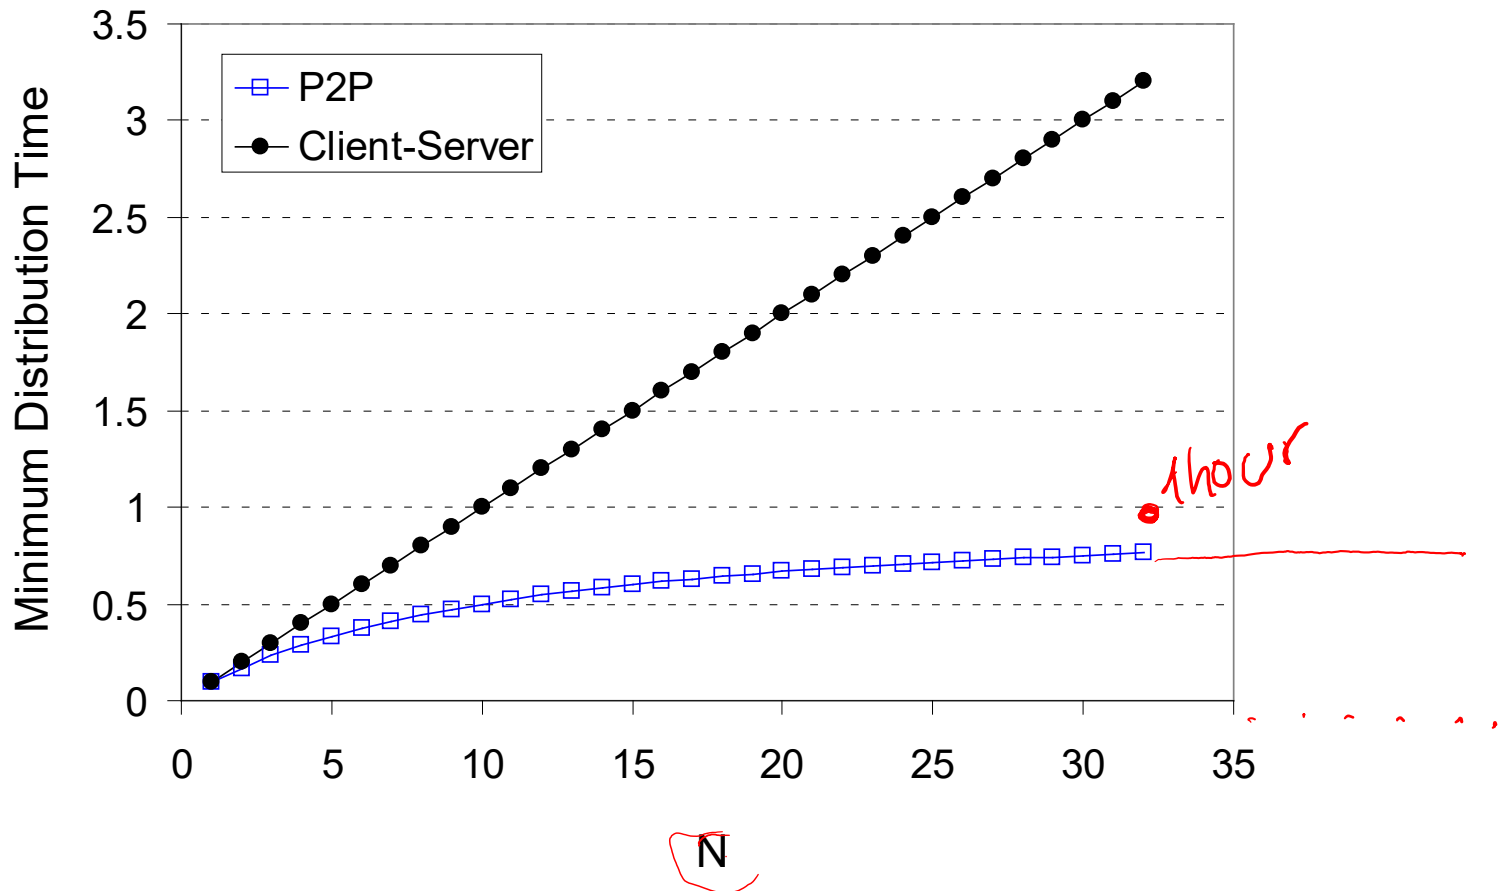
$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



Application Layer – Part 3


- ✓ P2P Applications
 - ✓ Characteristics
 - ✓ File distribution with BitTorrent
- ✓ Video streaming and Content Delivery Networks
 - ✓ Internet video
 - ✓ HTTP Streaming
 - ✓ CDN

P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent


torrent: group of peers exchanging chunks of a file



Alice arrives ...
... obtains list
of peers from tracker
... and begins exchanging
file chunks with peers in torrent

P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn**: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

TCP

requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Application Layer – Part 3

- ✓ P2P Applications
 - ✓ Characteristics
 - ✓ File distribution with BitTorrent
- ✓ Video streaming and Content Delivery Networks
 - ✓ Internet video
 - ✓ HTTP Streaming
 - ✓ CDN

Video Streaming and CDNs: context

- video traffic: major consumer of Internet bandwidth
 - Forecasted as 81% of consumer Internet traffic in 2021 (Source: Cisco)
 - Mobile Internet video traffic (usage measured from smartphones) accounts for a 49%, as of May 2021 (source: Visual Capitalist)
 - YouTube 48%, TikTok 16%, FB video 15%
 - ~116M Disney subscribers, ~73M Netflix subscribers in North America (source: Forbes.com)



Video Streaming and **CDNs**: context

- challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **solution**: distributed, application-level infrastructure




Multimedia: video

- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec / 30 fps
- digital image: array of pixels
 - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)


low-quality 100 kbps
HD 3 Mbps, 4K video 10 Mbps

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i




frame $i+1$

Multimedia: video

- **CBR: (constant bit rate):**
video encoding rate fixed
- **VBR: (variable bit rate):**
video encoding rate changes
as amount of spatial,
temporal coding changes
- **examples:**
 - MPEG I (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, < 1 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i

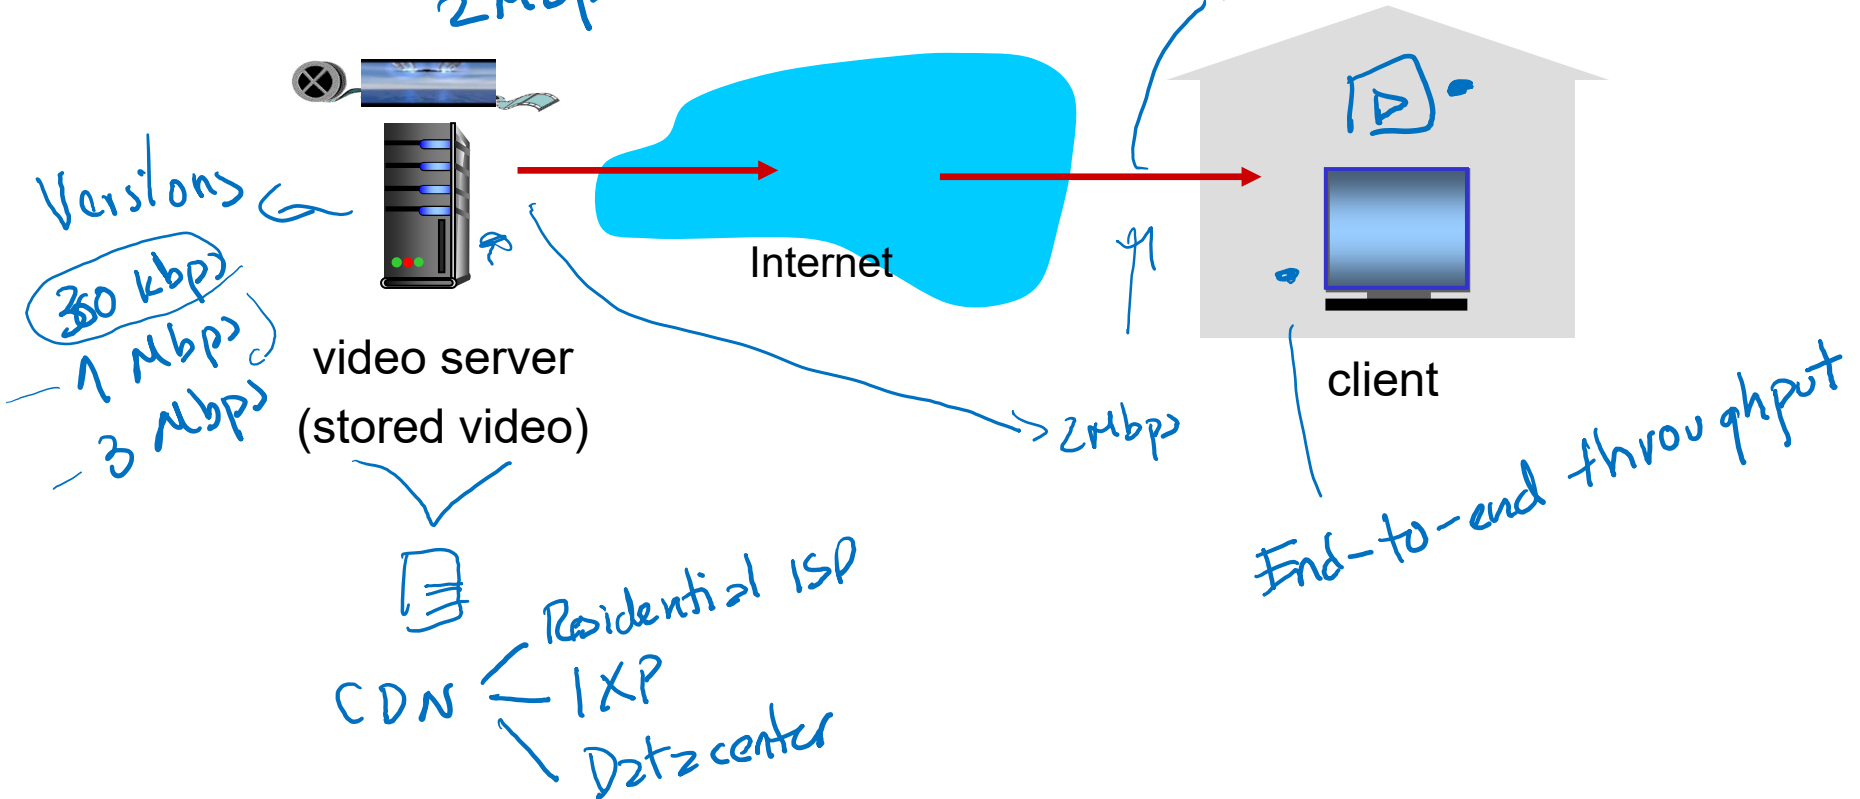


frame $i+1$

Streaming stored video:

simple scenario:

2 Mbps — 67 minutes, 1 GB



Application Layer – Part 3


- ✓ P2P Applications
 - ✓ Characteristics
 - ✓ File distribution with BitTorrent
- ✓ Video streaming and Content Delivery Networks
 - ✓ Internet video
 - ✓ HTTP Streaming
 - ✓ CDN

Streaming multimedia: DASH

- **DASH**: **D**ynamic, **A**daptive **S**treaming over **HTTP**
- **server**:
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
 - *manifest file*: provides URLs for different chunks
- **client**:
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- “intelligence” at client: client determines
 - when to request chunk (so that buffer starvation, or overflow does not occur)
 - what encoding rate to request (higher quality when more bandwidth available)
 - where to request chunk (can request from URL server that is “close” to client or has high available bandwidth)



End-end throughput

Application Layer – Part 3


- ✓ P2P Applications
 - ✓ Characteristics
 - ✓ File distribution with BitTorrent
- ✓ Video streaming and Content Delivery Networks
 - ✓ Internet video
 - ✓ HTTP Streaming
 - ✓ CDN

Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- **option 1:** single, large “mega-server”


- single point of failure
- point of network congestion
- long path to distant clients
- multiple copies of video sent over outgoing link



....quite simply: this solution **doesn't scale**


Content distribution networks

- *challenge*: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- *option 2*: store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
 - *enter deep*: push CDN servers deep into many access networks
 - close to users
 - used by Akamai, 1700 locations
 - *bring home*: smaller number (10's) of larger clusters in POPs near (but not within) access networks
 - used by Limelight

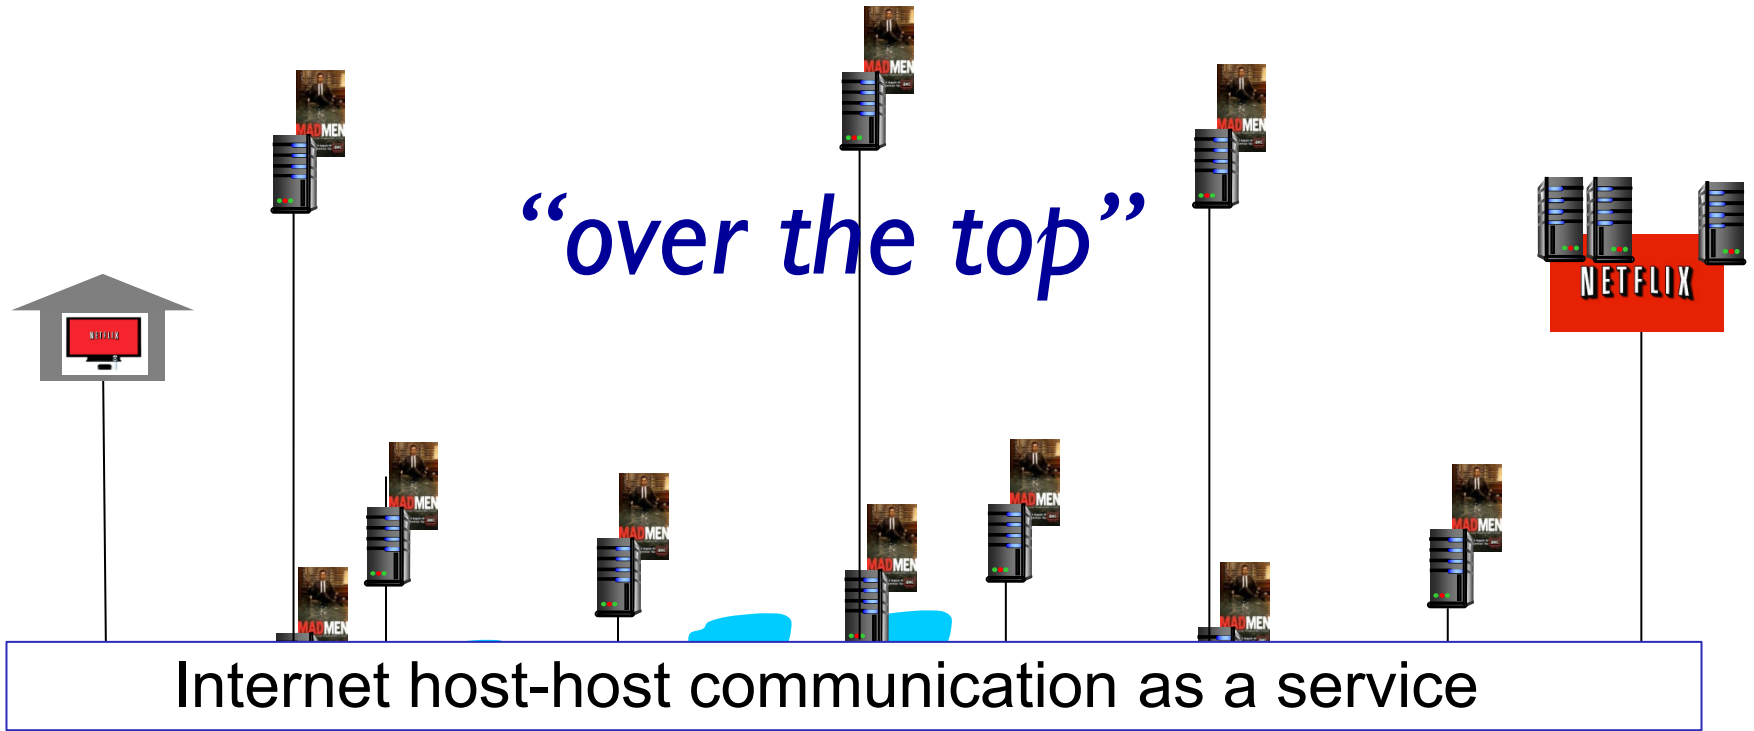


Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Content Distribution Networks (CDNs)



OTT challenges: coping with a congested Internet

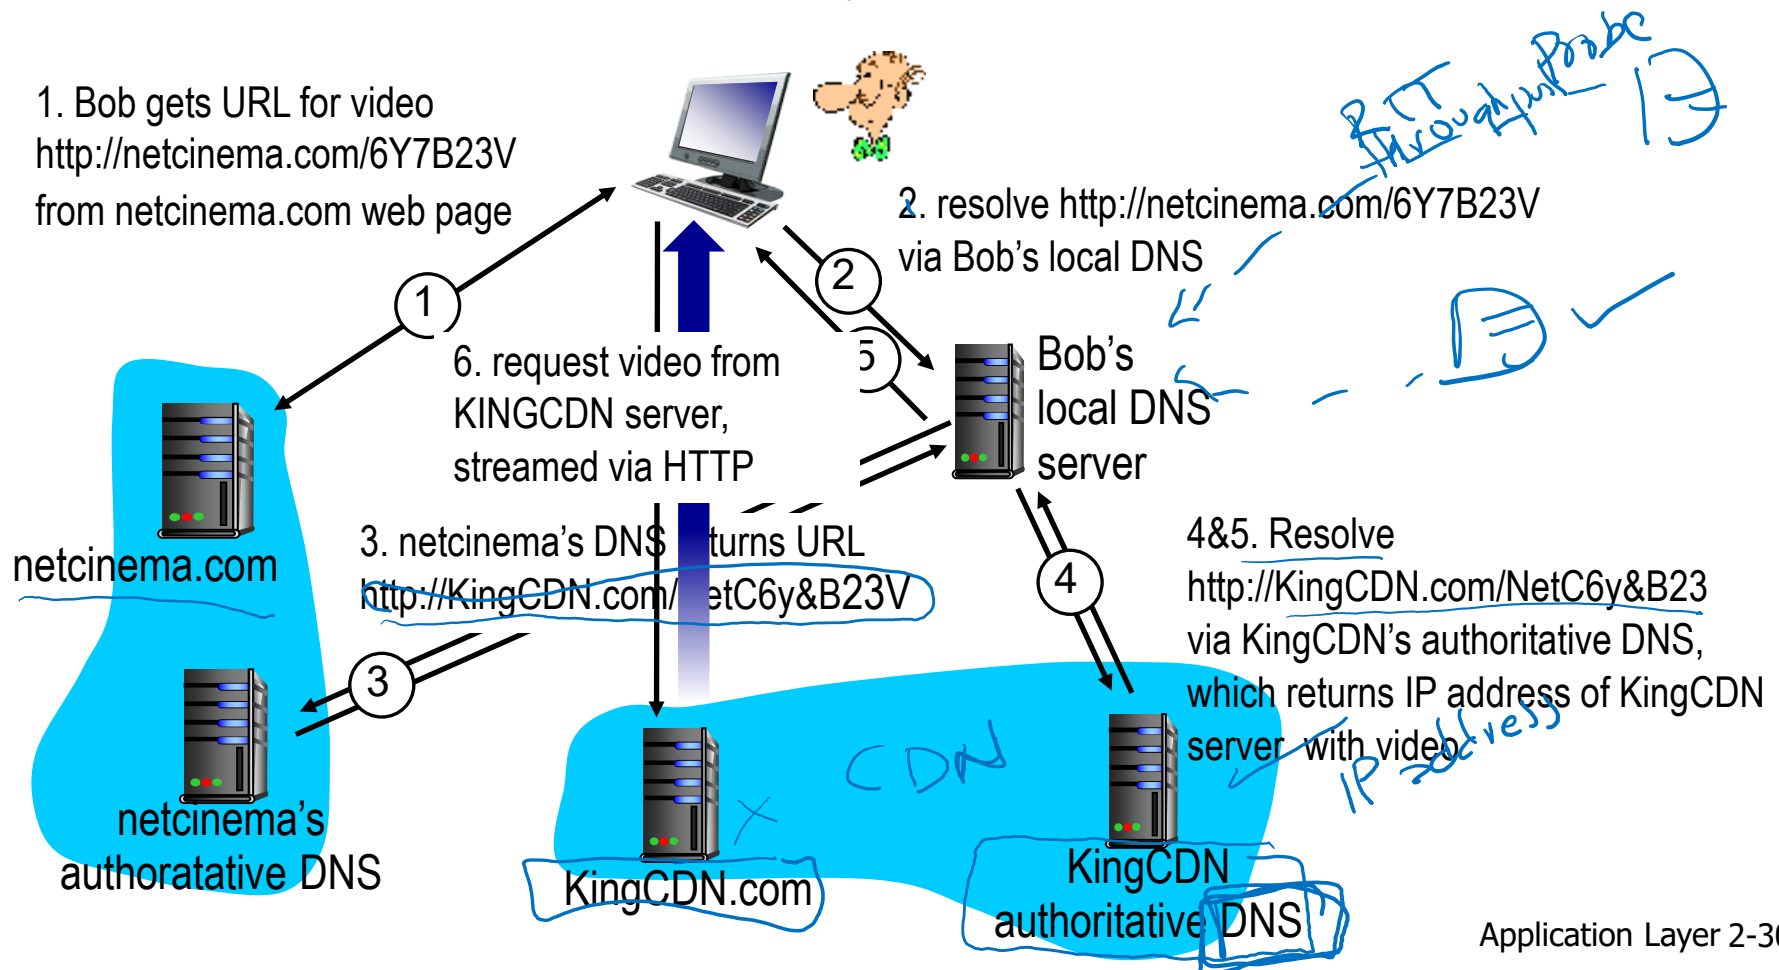
- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

more .. in chapter 7

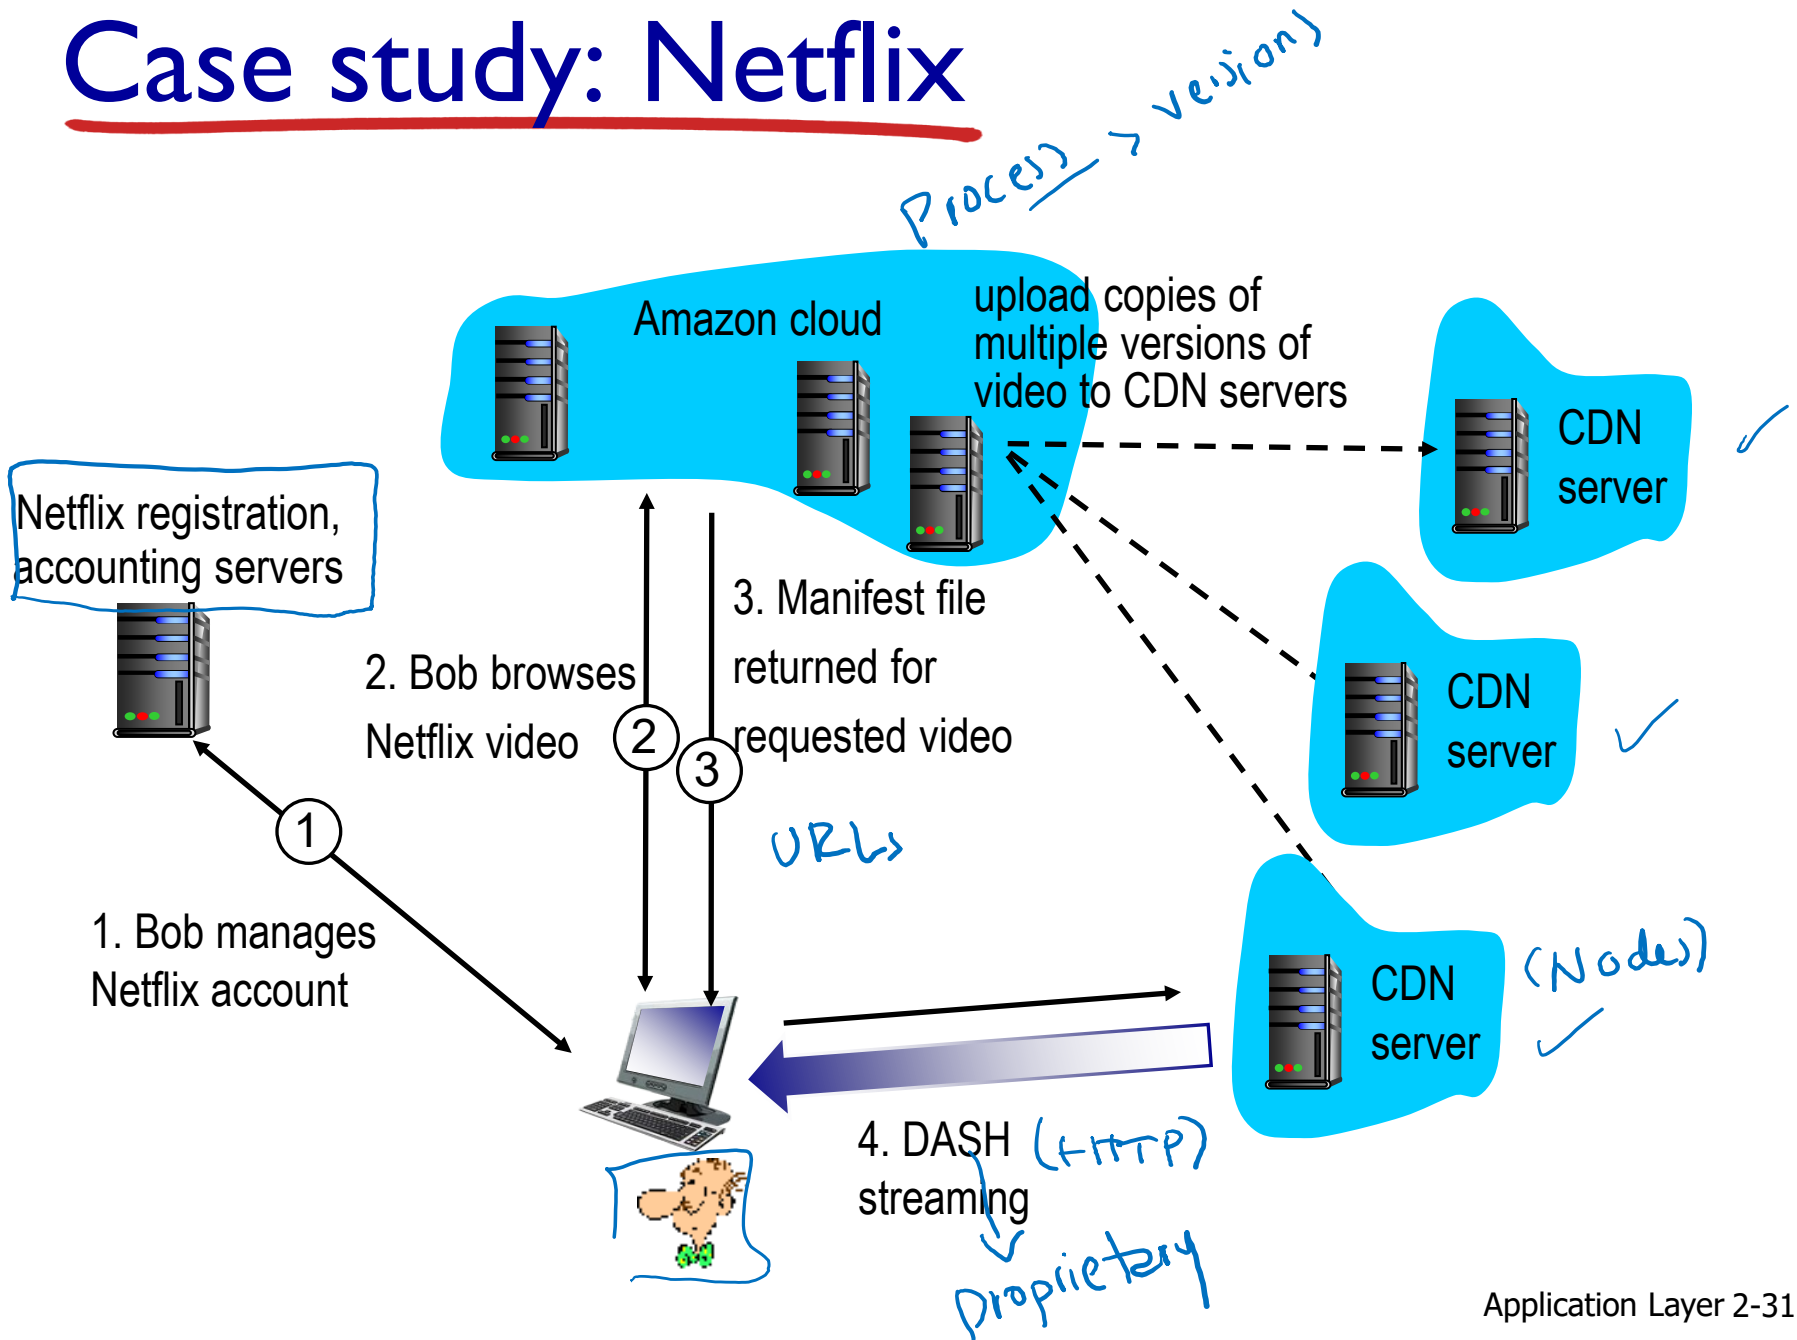
CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



Case study: Netflix



Case study: YouTube and Kankan

❖ *YouTube*

- Private CDN to distribute videos
- Uses pull-caching and DNS redirect
- Cluster selection based on measured RTT with load balancing
- Manual selection of video version

❖ *Kankan*

- Based on P2P *X client-server*
- Similar to BitTorrent
- Request for play-first chunks, to ensure smooth playback of videos

References

Figures and slides are taken/adapted from:

- Jim Kurose, Keith Ross, "Computer Networking: A Top-Down Approach", 7th ed. Addison-Wesley, 2012. All material copyright 1996-2016 J.F Kurose and K.W. Ross, All Rights Reserved